PRODUCTION, SCHEDULING AND CAPACITY PLAN

Running five YouTube channels without running out of anything

Three videos per week, per channel. Fifteen a week in total. This document sets out exactly when each channel writes its episode, exactly when each one publishes, how the free AI allowances are shared out so they never run dry, and eleven additional free services you are not yet using.

The short version

Everything below is explained in plain language further on. If you read nothing else, read these six points.

- **Three a week per channel is the right number.** It is the top of the band the growth research recommends, and it uses roughly a quarter of what your free accounts can supply in a day. It does not cost you reach.
- **Publishing time and writing time are two different problems.** Your audience is in America and Britain; their best hours fall in the middle of your night. The system already separates the two, so the machine can publish at 1 a.m. your time while you approve things at breakfast.
- **Each channel gets its own publishing hour**, because a finance viewer and a true-crime viewer do not watch at the same time of day. Finance goes out in the American morning; the documentary channels go out in the late afternoon.
- **Write the day before you publish.** This one change removes almost all the risk: if a night's writing fails, there is a whole day to fix it before the slot is missed.
- **Yesterday's problem was never volume.** Three separate faults let one run swallow a whole day's allowance. Two are fixed; the third is described on page 6.
- **There are at least eleven more free services available**, plus one source of unlimited free capacity that is not an outside service at all.

1.

Does three a week cost us reach?

No — and the evidence points the other way. This was your main worry, so it is worth answering properly rather than in passing.

What was studied	What was found	What it means for us
Upload frequency across many channels	One to three videos a week is the sweet spot for most channels. Beyond that, the algorithm rewards satisfaction and watch time, not volume.	Three a week is the **top** of the productive range, not a reduction.
A channel that went from 1 to 3 uploads a week	Total views rose 15%, but the average watch time per video fell 40%.	More videos can mean more views and a weaker channel at the same time.
A channel that slowed to one strong video a fortnight	Views per video rose 180%.	Quality per video is the stronger lever. That is what our scoring gates protect.
Publishing consistency	A steady schedule ("every Wednesday") beats bursts. Consistent channels get treated as higher priority.	Fixed days matter more than which days. Our plan fixes them.
Subscriber notifications	YouTube sends at most three notifications per channel per day.	Harmless at three a week. It would bite immediately if we went daily.

Where your growth actually comes from

Upload count is not the limiting factor, and pushing it higher is the one lever the research says can backfire. The levers that do move traffic are the click-through rate of the thumbnail, how long people stay, the topics chosen — and Shorts, which are the discovery engine that feeds subscribers to the long videos. **Shorts are currently switched off.** Turning them back on, once they are rebuilt, will do far more for reach than a fourth weekly upload ever would.

A word on the larger ambition

You have a very large revenue goal. I am not going to pretend a schedule decides that, and I would be doing you a disservice if I implied otherwise. What a schedule can do is make sure the machine runs every week without collapsing, that every video published is one you would defend, and that nothing is wasted. Fifteen well-made videos a week, published reliably for a year, is 780 assets. That is a real library. Fifteen a week that sometimes fails, publishes weak episodes and burns its own resources is not, however many were attempted.

2.

When each channel publishes, and why

Different subjects are watched at different hours. Treating all five channels the same would waste the advantage of having five. The rule underneath every row below: **publish two to three hours before the audience arrives**, so YouTube has time to index the video and start showing it to people just as the traffic builds.

Channel	Subject	Who watches, and when	Publish (UTC)	Your time
Collapse Index	Personal finance, retirement, startup failures	Working people, at their desks. Business and money content peaks in the American morning and at lunchtime, around the market open.	**11:00**	16:30 IST
No Known Cause	Medical mystery cases	Curious general audience plus medical students. Educational documentary — watched in the afternoon and evening, not first thing.	**19:00**	00:30 IST
Archive	Ancient civilisations, fallen empires	Long-form history is wind-down viewing: late afternoon onward, and strong at weekends when people have a spare hour.	**19:45**	01:15 IST
Control Files	Cults, propaganda, scams	Dark documentary behaves like true crime: evening, with a long late-night tail.	**20:30**	02:00 IST
Evidence Room	Forensic finance, criminal investigation	True crime proper — the latest-peaking of the five, and the strongest at weekends.	**21:00**	02:30 IST

What those times are in the audience's own clock:

Publish (UTC)	US East Coast	US West Coast	United Kingdom
11:00	6:00 am	3:00 am	12:00 noon
19:00	3:00 pm	12:00 noon	8:00 pm
19:45	3:45 pm	12:45 pm	8:45 pm
20:30	4:30 pm	1:30 pm	9:30 pm
21:00	5:00 pm	2:00 pm	10:00 pm

Why Collapse Index breaks the pattern

It is the only channel aimed at people at work rather than people relaxing. Money and business content is consumed on the commute, before the market opens and over lunch — not at 9 p.m. Publishing it at 11:00 UTC puts it in front of the American morning and the British lunch hour at the same time. Publishing it in the evening alongside the documentaries would waste it.

None of this needs you awake. Publishing is fully automatic and happens while you are asleep. The only thing that needs you is approving the episode, and that happens the previous morning — see the next section.

3.

When each channel writes its episode

Here is the single most useful change in this document.

Write the day before you publish, not the same day

At the moment an episode is written and published on the same day. If the writing fails — as it did three times yesterday — the slot is simply missed. If instead every episode is written the day before its slot, a failure costs you nothing: there is a full day to run it again. It also means the machine can write at whatever hour is quietest for the free allowances, independently of when the video needs to appear.

The writing week

Each channel writes three times a week, on fixed days, spaced two to three days apart. No more than three channels write on any one day, and never two at the same moment.

Channel	Mon	Tue	Wed	Thu	Fri	Sat	Sun
Collapse Index	A	—	A	—	A	—	—
No Known Cause	—	A	—	A	—	A	—
Archive	—	—	B	—	B	—	A
Control Files	B	—	—	B	—	B	—
Evidence Room	—	B	—	—	C	—	B

Writing slot	Starts (UTC)	Your time (IST)	Why this hour
A	03:30	09:00 am	Approval messages reach you over breakfast.
B	06:30	12:00 noon	Three hours after A, so the two never overlap and never compete for the same allowance.
C	09:30	03:00 pm	Friday only — the one day three channels write.

So a typical episode's life looks like this:

Time	What happens	Do you need to be there?
Tue 05:40 UTC	The health check tests every AI service and reports.	No. It only messages you if something needs replacing.
Tue 09:00 IST	Writing starts for Wednesday's No Known Cause episode.	No.
Tue 10:30 IST	Script ready. Telegram asks you to approve it.	**Yes — a few minutes**
Tue 12:00 IST	Voice and video built. Telegram asks again.	**Yes — a few minutes**
Tue 14:00 IST	Finished episode is stored, waiting.	No.
Wed 00:30 IST	It publishes automatically at the best hour for its audience.	No. You are asleep.

If you are away or asleep when an approval arrives

The episode is held, not published. A gate that nobody answered is recorded as "never reached you", never as your approval — that is deliberate, and it is why nothing goes out unseen. Because writing now happens a day early, a missed approval costs you a few hours, not the slot.

4.

How the free allowances are managed

This is the part you asked about most directly: with three or four channels producing on the same day, how does it never run out? Seven mechanisms, in order of how much they contribute.

First, the arithmetic

Measure	Number	Note
AI requests a healthy episode needs	30 to 90	Counted from the code. A troubled one can reach 400–600, which is what must be capped.
Episodes per day, on this plan	2, sometimes 3	15 a week across five channels.
Requests needed per day	roughly 130–190	
Requests the free accounts supply per day	roughly 700 today	Closer to 1,500 once the additions in section 6 are in.
How much of the ceiling we use	**about 20–25%**	Four to five times more headroom than the plan needs.

Second, the seven mechanisms

1. One channel at a time, never two

Two channels writing simultaneously ask the same services in the same minute and trip the per-minute limits, which looks exactly like an outage. Slots are three hours apart and the system refuses to start a second run while one is going.

2. Each channel starts at a different service

Rather than every channel hammering the same provider first, each begins its list at a different one and rotates. Five channels spread across the providers means no single allowance carries the week.

3. Small jobs go to small accounts, big jobs to big accounts

Writing a script needs a large allowance. Writing a six-word thumbnail line does not. The services with tiny daily limits but generous speed handle the short jobs; the ones with big daily token allowances handle the scripts. Sending a six-word request to the biggest account is pure waste.

4. Every request asks for only what it needs

Some requests currently ask for room for 8,000 words of output when the answer will be twenty. Measuring the question and asking for only the space the answer needs is the single biggest efficiency gain available, and costs nothing in quality.

5. A hard spending limit per episode

Every run counts its own requests. If an episode exceeds its budget it stops and skips the day, reported honestly as a capacity stop. One runaway episode can no longer eat the day — which is exactly what happened yesterday.

6. The morning health check is a gate, not a report

It runs before the day's episodes. If fewer than three services are healthy, the day's runs cancel themselves rather than discovering it the expensive way. Ten tiny test messages, versus five hours of failure.

7. Shared research, done once

The daily trend and research step produces the same answer for every channel on a given day. Doing it once and sharing it, instead of five times, removes a fifth of the day's requests for free.

5.

What actually exhausted the allowances yesterday

Not volume. Three specific faults. Two are now fixed and pushed; the third is the largest and is still to do.

Fault	What it was	Status
Retired models asked for over and over	When a service said "that model no longer exists", nothing wrote it down, so the same dead name was requested thousands of times. One run spent 117 minutes doing this and produced nothing.	**Fixed**
No stop signal when everything was down	With every service failing, the machine kept working through its thirty-nine attempts, then blamed the writing quality. It now stops after eight failed sweeps and says plainly that the services were down.	**Fixed**
Requests too large for the account that receives them	Our largest allowance is asked for more room than its free tier permits, so it refuses work it could otherwise do. One provider already measures the question and asks for the right amount; the others do not.	**To do — do this first**

An honest note about the numbers in this document

While researching this I found that published limits for these free services disagree between sources — in one case a provider's context limit is quoted as 8,192 by one source and 64,000 by another, and Gemini's daily allowance is quoted anywhere between 20 and 1,500 requests. I am not going to pretend to a precision I do not have. The fix works either way, and the right answer is not to trust any published figure: the daily health check should record what each service *actually* allowed, so that within a week we are working from our own measurements rather than somebody's blog post. That is standard 8.

6.

Eleven more free services you are not using

You were right that ten is not enough. These are all genuinely free, none requires a credit card unless noted, and all but two speak the same protocol as the services already connected — meaning they are a small amount of work each, not a rebuild.

Add these first — biggest gain, least effort

Service	What you get free	Why it is worth adding
Kilo Code	About 200 requests per hour	By far the largest free allowance on this list. On its own it roughly doubles today's daily ceiling.
Hugging Face	100,000 credits a month	Thousands of models including the Llama and Qwen families. Well documented and stable.
LLM7	30 requests a minute, 120 with a free token	No signup needed to start. Includes DeepSeek and GPT-class models.
ModelScope	2,000 requests a day	The largest daily request allowance here. Needs an Alibaba Cloud account, which is free.
SiliconFlow	30 requests a minute	Qwen and DeepSeek models, generous speed limit.

Add these second — useful, smaller, or slightly more setup

Service	What you get free	Note
OVHcloud AI	2 requests a minute, per model	**No signup at all.** Anonymous access. Slow, but it can never expire or be revoked — an ideal last resort.
Z.AI (GLM)	Free, one request at a time	Chinese lab, strong models, no daily cap published.
Aion Labs	15 requests a minute, 20,000 words a day	Small but genuinely permanent.
Ollama Cloud	400+ models, weekly session limits	Speaks a different protocol, so slightly more work to connect.
Cloudflare Workers AI	10,000 units a day	Already connected, but currently only three models are wired up out of fifty. Widening this is nearly free.
GitHub Models	10 requests a minute, 50 a day	We retired this after it returned "gone". Current listings show it alive with 45+ models, so our connection details are probably just out of date. Worth one retest.

And one that is not an outside service at all

Run a small AI model on the build machine itself — unlimited, free, unkillable

Your repository is public, which means GitHub gives you build machines free and **without any monthly limit**. Each one has 4 processors and 16 GB of memory. A small AI model can be downloaded and run directly on that machine.

It would be too slow for writing a full script. But roughly half of all the requests an episode makes are tiny — a title, a six-word thumbnail line, a set of hashtags, a yes/no judgement. For those, a local model answers in seconds.

This matters more than any single service above, because it has **no key, no quota, no rate limit, and nobody can decommission it.** It is the one provider that can never appear in a morning report as dead. Every problem we have had for a month — expired keys, retired models, exhausted allowances — is structurally impossible for it.

Honest caveat: I have not benchmarked the speed on that exact hardware, so this needs a half-day test before it is relied upon. I am proposing it, not claiming it works.

7.

Eight standing rules

Written so this conversation does not need to happen again. Each is numbered so it can be referred to, and each will be enforced by an automatic check — so a future change that breaks one fails a test rather than a week of uploads.

Rule 1 — Every run has a spending limit

Each episode counts its own AI requests and stops if it exceeds its budget. A run that cannot say what it spent cannot be planned around.

Rule 2 — One channel produces at a time

Minimum three hours between starts, enforced automatically. Never two at once.

Rule 3 — The morning health check decides whether the day runs

Fewer than three healthy services and the day cancels itself, before spending anything.

Rule 4 — Every request asks only for the room it needs

Measure the question, subtract from what the service allows, request the remainder. No request inherits a default it cannot use.

Rule 5 — Writing happens while you are awake; publishing does not

Anything needing your approval runs between 09:00 and 15:00 your time. Anything automatic runs at the audience's best hour. These two never merge.

Rule 6 — Write one day ahead of publishing

Every episode is finished the day before its slot, so a failure has a day of recovery instead of costing the upload.

Rule 7 — One-off credits are never counted as regular capacity

Some services give a fixed sum that never renews. These are labelled as expiring and excluded from the daily total, so nothing important depends on them.

Rule 8 — We measure the limits ourselves rather than trusting the published ones

The health check records what each service actually permitted, building our own table over time. Published figures change without notice and disagree with each other.

What I would build, in this order

	Work	Why here
1	Make every request ask for the right amount of room	Unlocks the largest allowance you already have. Costs nothing, gains the most.
2	Add the request counter and the per-episode limit	You cannot manage what you do not measure — and it makes rule 1 enforceable.
3	Get one clean episode through end to end	Nothing else is trustworthy until one full episode completes. None has since yesterday's fixes.
4	Connect the five easiest new services	Roughly doubles the daily ceiling for a few hours of work.
5	Set the writing and publishing timetable	Quick to do, but pointless before step 3 succeeds.

	Work	Why here
6	Test the local model on the build machine	Half a day to find out whether an unlimited free provider is available to you.
7	Bring channels 2 to 5 up to channel 1's standard	The biggest job. Every fault fixed on channel 1 this month probably exists in the other four, and none of them has ever run.

The one thing to hold me to

Channel 1 has not completed a successful run since these fixes went in. Every number in this document about how much an episode costs is calculated from the code, not measured from a finished episode. Rule 1 exists precisely to replace those estimates with real measurements. Expect this plan to be adjusted once, after the first clean week — and treat anyone who tells you a plan like this is final, before a single run has finished, with suspicion.